
BitTensor: An Intermodel Intelligence Measure

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 The dominant approach to creating a machine learning system is to collect a dataset
2 of training examples that demonstrate correct behavior for a desired task, then
3 train that system to imitate these behaviors [1]. Labeled datasets that evaluate the
4 performance of intelligent systems on specific predefined tasks – also known as
5 benchmarks – are the typical tools used to measure and reward the advancement
6 of machine intelligence. We argue that a reward mechanism based on these tasks
7 where progress is measured as test scores is low resolution, misaligned, and in-
8 troduces wide-scale inefficiencies in developing machine learning systems. To
9 reconcile these issues, we propose BitTensor: a benchmark where intelligence is
10 measured by other intelligence systems rather than against task-based tests. We
11 describe how this novel benchmark is constructed and how it is negotiated by com-
12 puters that share knowledge peer-to-peer (P2P) across the internet. However since
13 negotiation introduces game-theoretic considerations, we design methods to ensure
14 veracity of the benchmark when peers remain self-interested and trustless. To test
15 our design, we empirically show that the self-interested negotiation produces the
16 desired ranking at the competitive equilibrium. The result is a benchmark which
17 is simultaneously higher dimensional and divisible allows for a much larger and
18 continually expanding definition of intelligence is inherently collaborative, open,
19 decentralized, and can consume resources connected by an internet fabric.

20 The best in-class machine learning systems today boast high performance and trend-setting bench-
21 marks at specific tasks for which they are trained. They utilize representations of inputs that are
22 semantic, disjoint, and general enough to tackle a given task while maintaining their integrity [1, 2].
23 However, when dealing with unexpected tasks or slight changes in data distribution these models
24 become brittle and sensitive, making them narrow experts rather than resilient generalists [1]. This is
25 due to the models’ reliance on training feedback mechanisms that benchmark them against examples
26 of data rather than direct knowledge [3].

27 Machine learning systems are trained on a dataset of training examples that demonstrate desired
28 behavior for a given task, and then its intelligence (i.e. performance) is then evaluated on independent
29 and identically distributed examples [1, 4]. Finally, the system’s performance on that task is measured
30 against existing systems as a “performance benchmark”. In contrast, if the mechanism is misaligned
31 to the salient feature being produced, then a large portion of the work is lost as a result. This low
32 resolution and sensitivity to changes in data is further compounded when considering the following:

- 33 1. Lack of breadth: the narrowness of the domain described by the tasks make it easier for
34 models to “buy progress” without improved general understanding [4].
- 35 2. Lack of inter-model knowledge transfer: measuring how much one model could improve
36 another is not possible due to the inherent incompatibility between model architectures.

37 After training, these models become very good at performing only that desired task. We argue that the
38 measure of intelligence should be – in information-theoretic terms – with respect to *intelligence* itself

rather than via a projection onto labeled datasets. Further, the lack of inter-model knowledge transfer ensures that the field is inherently non-collaborative, with the majority of researchers world-wide in essence contending against each other [5]. Ultimately, the only users capable of creating the systems with high performance are those with the means and resources to train the largest models.

We present a novel framework in which a collective of intelligence systems can be used to evaluate their peers. By using intelligence itself as the measure of a model’s intelligence, this allows models to share knowledge with each other and removes the need to re-train models when new information arrives. Finally, this turns the performance benchmark into an intelligence market that rewards models that improve knowledge within it.¹ We design the benchmark to run in a continuous and asynchronous fashion across a peer-to-peer (P2P) network. Since this introduces trustless computation² we also dedicate a part of this paper to explaining how the system remains fair when little assurance can be given about the computers that compose it.

The contributions of this paper are thus four-fold:

1. We introduce a P2P based intelligence framework.
2. We introduce an “intelligence market” that rewards models that increase the knowledge within it.
3. We show how models can share knowledge to increase each other’s performance and remove the need to re-train models each time new data arrives.
4. We propose turning typical data-based performance benchmarks into an intelligence market, in which models that improve knowledge are rewarded.

The rest of this paper is organized as follows: Section 1 presents the intelligence benchmark and ranking mechanism that models use to evaluate each other. Section 2 considers the scenario where participants are not honestly reporting the significance of their peers, and proposes a method to detect such behavior. Section 3 highlights the experiments performed to verify the function of the framework. Finally, Section 4 presents a summary and future works on this system.

1 Methodology

1.1 Benchmark

The network is composed by n unique parameterized functions $F = f_0, \dots, f_j, \dots, f_n$ where each function is producing an output tensor $f_i(F(x))$, a “representation” from an input tensor $F(x) = [f_0(x) \dots f_n(x)]$ gathered by querying its neighbors. Each function is training asynchronously over a dataset $D_i = [X, Y]$ such that, given an error function Q_i , its expectation over that data E_{D_i} defines a loss $\mathcal{L}_i = E_{D_i}[Q_i(y, f_i(F(x)))]$. We assume these losses are measured on the same scale and thus our benchmark \mathcal{B} can be defined by their sum:

$$\mathcal{B} = \sum_i^n \mathcal{L}_i \quad (1)$$

¹“The iron rule of nature is: you get what you reward for. If you want ants to come, you put sugar on the floor.” - Charlie Munger

²In decentralized systems, trust is shifted from the individuals to the network itself, “trustless” usually means “minimal trust is required” w.r.t the protocol

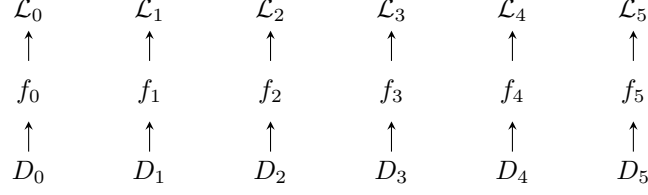


Figure 1: $n = 6$ parameterized functions with losses \mathcal{L}_i and datasets D_i .

Each parameterized function is represented here in its most abstract sense[6] and need only accept the same input type x and produce the same output dimension to fit within the network. For instance, unicode encoded string $x = \text{"hello"}$ and its semantic representation as a 64×49 word embedding. This widened scope ensures participants can be multi-task [7], use completely distinct computing substrates [8] or train on unique datasets. [9].

1.2 Ideal Ranking

Our goal in this work is to produce a ranking $R = [R_i]$ over these functions where the score $R_i \in R$ represents participant i 's information-theoretic significance to the benchmark \mathcal{B} . Following Le Cun and others [10, 11], it is reasonable to analytically define this significance by equating it with the cost of removing each component from the network:

$$R_i \approx \frac{1}{n} \sum_j \sum_{x \in D_j} \Delta F^T(x)_i * H(\mathcal{Q}_j(x)) * \Delta F(x)_i \quad (2)$$

$$\Delta F(x)_i = [0, \dots, 0, -f_i(x), 0, \dots, 0]$$

Where the above is derived using a Taylor series (Appendix 6.1) and $\Delta F(x)_i$ is the perturbation of the i^{th} node's inputs as a function of it's choice of weights when removing function f_i at the point x . Note, the linear and higher order terms of the Taylor series have been removed following [11] and the remaining term $H(\mathcal{Q}_i)$ is the hessian of our error function. When the error function \mathcal{Q} is the twice-differentiable cross-entropy, then $H(\mathcal{Q}_i)$ is the Fisher information matrix, and $R_i \in R$ is measured as relative entropy: reflects each participants informational significance to the network as a whole.

1.3 Inter Ranking

It is not possible to compute the ranking score above without access to the parameters of each function in the network. Instead, we use a set of inter-model weights $W = [w_{i,j}]$ where each $w_{i,j}$ is the score attributed to f_j from f_i combined into an $n \times n$ square matrix.

$$w_{ij} = \sum_{x \in D_i} \Delta F^T(x)_j * H(\mathcal{Q}_i(x)) * \Delta F(x)_j \quad (3)$$

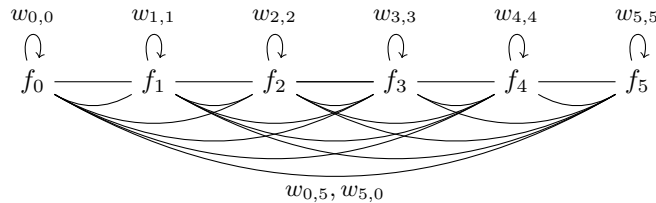


Figure 2: Inter-model contribution weights: $w_{i,j}$ the score attributed to f_j from f_i

93 The weights can be computed on the fly either approximately [11] or by using the full hessian of
 94 the error. We store them on a distributed ledger and allow participants to update them by making
 95 changes of bounded size: $W^{t+1} = W^t + \lambda \Delta W$, where $\|W_i\|_2 < \epsilon$ at block step t . We also enforce
 96 that the scores in each row sum to 1, $\|w\|_1 = 1$. **The reasoning for this enforcement is to ensure**
 97 **stake fairness between early movers and newcomers with innovations that enable them to perform**
 98 **better overall.**

99 The equivalent ranking R (2) can then be computed by normalizing column sum of the weight matrix:
 100

$$R = \frac{1}{n} W^T * \mathbb{1} \quad (4)$$

101 The problem is that without system-wide access to the model parameters, the computation of w_{ij} (3)
 102 is intractable, **as it would require collecting great amounts of information from potentially billions of**
 103 **neurons.** It is reasonable to assume participants will select weights which artificially increase their
 104 own rank rather than others in the network. Moreover, since the network remains open, participants
 105 may choose to create many spuriously neighbours and rank themselves higher. The remainder of this
 106 paper describes our proposal for resolving these issues.

107 1.4 Stake

108 The proposed solution begins by introducing a finite resource $S = [s_i]$, a component's 'stake' in the
 109 system, and an inflation mechanism τ which translates the ranking vector R into additional stake as
 110 incentive.

$$R = \frac{1}{n} W^T \circ S * \mathbb{1} \quad (5)$$

$$S^{t+1} = S^t + \tau * \frac{R}{\|R\|_2} \quad (6)$$

111 The \circ here is the Hadamard product between the $n \times n$ matrix W and the $n \times n$ matrix containing S
 112 in each column, and t is the time-step referred to in Section 1.3 (measured in distinct blocks on the
 113 distributed ledger). By design (5) increases the importance of those with stake, $s_0 * w_{ij}$. This serves
 114 two purposes:

- 115 1. New computers can spuriously create new nodes, however this won't game the ranking
 116 because the amount of stake they hold is finite.
- 117 2. The resource provides mechanism power.

118 By providing it to nodes with large rank, this ensures that those with weight must have worked to
 119 attain it, or indirectly subsidized those who have done so already. A single staked token would be
 120 enough to bootstrap the process.³

121 1.5 Competitive weights

122 While stake provides some protection against malicious actors, it does not ensure weights are set
 123 accurately. Our solution begins by introducing competition for connectivity within the network.
 124 Nodes that underweight are punished by having inputs from the network masked to zero (7), **which**
 125 **in turn would punish the loss function.** To frame this market we borrow the continuous differential
 126 activation function σ with range $(0, 1)$. Under a choice of weights W_i the inputs to component i are:

³The propagation of a resource mocks the flow of the brain-derived neurotrophic factor (BDNF). [12]

Algorithm 1 Inflation mechanism

Require: $S = [n \times 1]$ Stake Vector
Require: $W = [n \times n]$ Weight Matrix
Require: $\tau > 0$ inflation rate
 while TRUE **do**
 $W = W + \lambda \Delta W$
 $R = \frac{1}{n} W^T \circ S * \mathbb{1}$
 $S = S + \tau * \frac{R}{\|R\|_2}$
 end while

$$F_W(x) = [f_0(x) * \sigma(s_i * w_{i,0} - \mu_0), \dots, f_n(x) * \sigma(s_i * w_{i,n} - \mu_n)] \quad (7)$$

$$\sigma = \frac{1}{1 + e^{-\frac{x}{T}}} \quad (8)$$

127 Here, the shift term μ_j is the average of the weights in each column $\mu_j = (\frac{1}{n}) \sum_i s_i * w_{i,j}$, and the
128 activation function is the temperature scaled sigmoid. Because the allocation mechanism is standard
129 across the network it is possible for each participants to compute both $\frac{\partial \mathcal{L}_i}{\partial W_i}$ and $\frac{\partial R_i}{\partial W_i}$. Computers may
130 augment their usual training framework, for instance, Tensorflow, with the allocation mechanism
131 shown here.

1.6 Running the network

133 The steps to run a network participant are:

- 134 1. Participant defines its dataset D_i , loss \mathcal{L}_i and parameterized function f_i
- 135 2. At each training iteration the participant broadcasts batches of examples from D_i to its peers
136 $[batch_size, x]$
- 137 3. Responses $F(x)$ from the network produce a loss-gradient $\frac{\partial \mathcal{L}}{\partial F}$ which back-propagates
138 through f_i and out to the network.
- 139 4. During 2 and 3 the participant competitively selects the weights for their row $w_{ij} \in W$.
- 140 5. Participants submit changes to the weights ΔW_i which, in-turn changes the ranking and
141 induces inflation $\tau * R$.
- 142 6. [After steps 1-5 have been performed multiple times](#), participants disconnect and verify the
143 model in a normal manner.

144 Peers only communicate with computers that hold stake as a consequence of section ???. Those that
145 fail to produce value will be pruned naturally as participants learn to differentiate signal from noise.

1.7 Conditional computation

147 As the network grows, outward bandwidth will become the major bottleneck. Components learn to
148 trim outward bandwidth by employing a Sparsely-Gated Mixture-of-Experts (SGMoE) [13] layer at
149 the input. The gating layer determines a sparse combination of children to query for each example
150 and then re-joins them using the the gating weights $g_j(x)$. The combined gated inputs are fed as
151 input to the local function:

$$f_i = f_i(G(x)) \quad (9)$$

$$G(x) = [\dots, g_j(x) * f_j(x), \dots] \quad (10)$$

152 The layer cuts outward bandwidth, querying only a small subset of peers for each example. The
 153 gating function is trainable w.r.t to the loss and its weights act as a proxy for importance $w_{ij} \in W$.
 154 This method has been shown to drastically increase the potential for outward bandwidth in datacenter
 155 training,[13] and has been investigated in a peer-to-peer (P2P) setting as well [5]

156 1.8 Extracting knowledge

157 Inter-node dependence in the network is broken using distillation[6], a compression and knowledge
 158 technique in which a smaller model – the student – mimics the behaviour of an ensemble. We
 159 employ this technique over the gating ensemble (10) where the student model learns to minimize
 160 the cross-entropy (shown below as KL) between the logits produced by the gating network and its
 161 predicted distribution. [14]

$$\text{distillation loss} = \text{KL}_D(\text{dist}(x), G(x)) \quad (11)$$

162 We use the distilled model as proxy to cut recursive calling between each components rather than
 163 query farther into the network. If models go offline, their peers can use their distilled versions in-place.
 164 Private data can be validated over the distilled models instead of querying the network. Eventually,
 165 components can fully disconnect from the network using the distilled inputs to validate and inference
 166 the models offline.




Figure 3: Queries propagate to depth=1 before the distilled model is used.

167 2 Analysis

168 We consider the scenario where participants are not honestly reporting the significance of their peers.
 169 The network is progressing at discrete timesteps t and it is reasonable to assume that each participant
 170 is attempting to maximize their subjective payoff over these steps.

171 2.1 Payoff Model

172 The staking system in Section 1.4 gives participants incentive to maximize their self-weight w_{ii} while
 173 the competitive connectivity described in Section 1.5 makes it costly for participants to decrease the
 174 remaining weights w_{ij} in their row. Since the row must sum to 1, we have a trade-off in two terms:

- 175 1. A utility term attached to the loss $U(\mathcal{L}(W))$.
- 176 2. The token emission via inflation $\tau * R_i(W)$.

177 Both of these are functions of the weights and, without loss of generality, measured in similar units:

$$P_i(W) = U_i(\mathcal{L}_i(W)) + \tau * R_i(W)_i \quad (12)$$

178

179 It is reasonable to assume payoff maximizing participants will use P_i as their objective during
 180 training. This can be computed using standard tools since both terms $U(\mathcal{L}_i(W))$ and $R(W)_i$ are fully
 181 continuous and differentiable. The system can be characterized as a competitive gradient descent
 182 game where participants are making steps $\Delta W_i = \frac{\partial P_i}{\partial W_i}$. In appendix 6.2 we prove that this strategy
 183 is regret-free and achieves the best expected payoff in hindsight. This assumption is also generally
 184 employed in smooth markets [15]. The iterative descent thus follows:

$$W^{t+1} = W^t + \lambda \Delta W \quad (13)$$

185

$$\Delta W = [\frac{\partial P_0}{\partial W_0}; \dots; \frac{\partial P_n}{\partial W_n}] \quad (14)$$

186

187 2.2 Empirical Model

188 Without access to the running network we evaluate our system using an empirical model. To derive
 189 the gradient steps in this model $\frac{\partial P_i}{\partial W_i}$ we make the following assumptions:

- 190 1. The utility functions are continuous-differentiable and are bounded by their first order
 191 derivatives $\frac{\partial U}{\partial \mathcal{L}} = \alpha$.
- 192 2. The network is converged to a local minimum in the inputs $\frac{\partial \mathcal{L}}{\partial F} = 0$.

193 The first assumption is approximate for small changes of continuous functions and the second
 194 assumption is realistic after extended training. In Appendix 6.2 we derive the following gradient step:
 195

$$\frac{\partial P}{\partial W} = \frac{\alpha}{\tau} * \frac{\partial L}{\partial W} + \frac{\partial R}{\partial W} \quad (15)$$

196

$$\frac{\partial L}{\partial W} = \frac{\partial}{\partial W} [(F_W - F_{W_0})^T * H(\mathcal{L}(F)) * (F_W - F_{W_0})] \quad (16)$$

197

198 Here F_W is the masked inputs from Section 1.5. $(F_W - F_{W_0})$ is the difference in the mask between
 199 the choice of weights W and the weights at the minimum W_0 and $H(\mathcal{L}(F))$ is the $[n \times n]$ hessian of
 200 the loss over inputs F . This formulation is both intuitive and useful: the gradient term $\frac{\partial L}{\partial W}$ measures
 201 the change in loss for any choice of weights. Additionally, it also allows us to compute the ranking
 202 for simulated hessian terms.

203 The remainder of the network is deterministic, and can be described by a choice of $\theta =$
 204 $[\alpha, \tau, \lambda, n, \sigma, T, S, W, H]$. For instance, the secondary term, $\frac{\partial R}{\partial W}$ can be computed-directly from
 205 Section 1.4 and only depends on the stake vector S and weights W .

206 3 Experiments

207 To generate sample statistics from the network, we first select $[\alpha, \tau, \lambda, S, \sigma, T, n] \in \theta$, then we
 208 generate random positive semi-definite Hessians H , and random uniform initial weights W_0 . For each
 209 parameterization we discover the competitive ranking by converging the system to a Nash-equilibrium:
 210 an equilibrium where no individual can vary from their set of weights and stand to gain. [16] To find
 211 these equilibrium, we use the competitive descent strategy described in (14) and compute the gradient
 212 terms from (15). In each trial we use a learning rate $\lambda = 0.005$ and stop when the gradient terms are

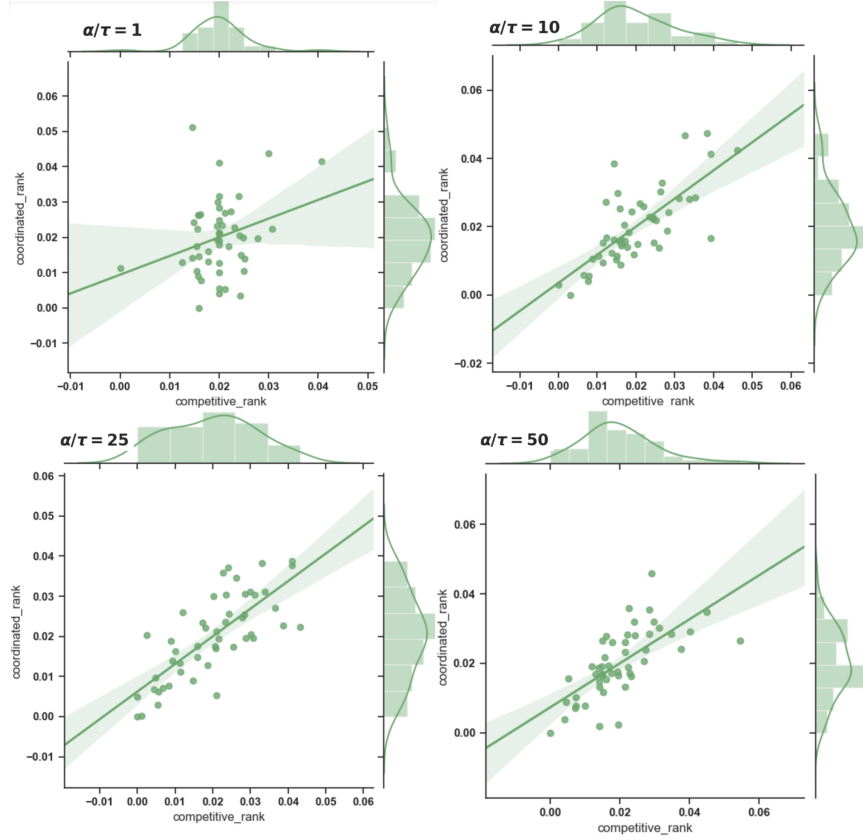


Figure 4: Correlations between the competitive rank and coordinated rank for $\frac{\alpha}{\tau} \in \{1, 10, 25, 50\}$. For low values of $\frac{\alpha}{\tau}$ the weights converge to the identity: the state where peers are fully disconnected.




Figure 5: (i) Left: the ratio between the main diagonal and the remainder of the weights. Right: the adaptive τ parameter converging onto the target. The weight matrix sparsity is a proxy for the ranking accuracy which we see in Figure 4. As sparsity converges onto the target the ranking is also improving.

213 bounded by ϵ or the steps exceed $1000 \times n$. The competitive ranking R^* at this point follows from
 214 (4) and can be compared to the idealized score in R from (2).

We are interested in the correlation between R^* and R as we vary the ratio α and τ , this ratio is explicit in (15) where the fraction relates the two gradient terms. Intuitively the ratio is between the value of minimizing the loss and maximizing revenue from inflation. Since this effects the ranking, we show this trade-off for various choices in Figure 4.

Finally, in Figure-5 we implement an adaptive- τ strategy where the network varies the inflation rate. Initial inflation is zero and then increases until the weights begins to converge towards the main diagonal $w_{ii} = 1$. We measure the sparsity $sparsity = \text{sum}(W_{dg})/\text{sum}(W)$, the ratio between the main diagonal and the remaining weights. As sparsity increases we push the market equilibrium by decreasing τ . Figure shows this adaptive convergence for $\alpha = 1$ with a sparsity target of 1.

3.1 Discussion

Figure 4 shows the relationship between the idealized rank and the competitive rank as a function of $\frac{\alpha}{\tau}$. Figure-4-a, $\frac{\alpha}{\tau} = 1$ shows the case where this ratio negatively effects the ranking accuracy. Here, all components have set $w_{ii} = 1$ and the resulting scores for all participants has converged to $1/n$. When this occurs, the system could decrease the inflation rate τ and push the network towards the high information markets seen in Figure-4-b, c, and d. Figure-5 shows a basic implementation of this where τ adapts to the ratio between the row-sums and the main-diagonal. By lowering inflation, it subsequently 'costs less' to connect with peers. Profit maximizing nodes automatically adjust to the change and the system converges back towards an accurate ranking. Those with high ranks will oppose inflation decreased, while those with low ranks will welcome it. The equilibrium found in this meta game will most certainly depend on the number of participants, key to both the ranking accuracy and the market at its core. However, we leave this analysis for a follow up paper.

4 Conclusion

We have proposed an inter-model benchmark that can run in a P2P setting outside of a trusted environment. We started with a typical machine learning framework defined by a set of functions with their losses over given datasets, then derived an idealized ranking score. The measure produced an information theoretic score that new participants can improve by learning how to be useful to their peers. However, the system is incomplete without a mechanism that prevents participants from ranking dishonestly. To resolve this, we proposed an incentive scheme and a differential allocation system over the network weights. The system allows the participants to train for connectivity in the graph. Following this, we described how to increase the outward bandwidth in the system using a trainable gating network and how to cut independence between nodes using distillation. Finally we showed how increasing the number of nodes in the system and fixing the inflation mechanism properly ensured that the resulting rank scores would correlate with those found in a idealized setting. While this is true, stake in the system holds value as a means to drive what the network learns. That benchmark is continually being solved by the participants, compounding what has been learned before and making it available to new learners in the system.

5 Broader Impact

In order to provide a balanced perspective, authors are required to include a statement of the potential broader impact of their work, including its ethical aspects and future societal consequences. Authors should take care to discuss both positive and negative outcomes.

References

- [1] A. Radford, J. Wu, R. Child, D. Luan, D. Amodei, and I. Sutskever, "Language models are unsupervised multitask learners," *OpenAI Blog*, vol. 1, no. 8, p. 9, 2019.
- [2] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "Bert: Pre-training of deep bidirectional transformers for language understanding," *arXiv preprint arXiv:1810.04805*, 2018.
- [3] A. Wang, A. Singh, J. Michael, F. Hill, O. Levy, and S. R. Bowman, "Glue: A multi-task benchmark and analysis platform for natural language understanding," *arXiv preprint arXiv:1804.07461*, 2018.
- [4] F. Chollet, "On the measure of intelligence," *arXiv preprint arXiv:1911.01547*, 2019.

- [5] M. Riabinin and A. Gusev, "Learning@home: Crowdsourced training of large neural networks using decentralized mixture-of-experts," *arXiv preprint arXiv:2002.04013*, 2020.
- [6] G. Hinton, O. Vinyals, and J. Dean, "Distilling the knowledge in a neural network," *arXiv preprint arXiv:1503.02531*, 2015.
- [7] L. Kaiser, A. N. Gomez, N. Shazeer, A. Vaswani, N. Parmar, L. Jones, and J. Uszkoreit, "One model to learn them all," 2017.
- [8] M. A. Nugent and T. W. Molter, "Cortical processing with thermodynamic-ram," 2014.
- [9] G. Lample and A. Conneau, "Cross-lingual language model pretraining," 2019.
- [10] Y. LeCun, D. J. S., and S. S. A., "Optimal brain damage," *Advances in Neural Information Processing Systems 2 (NIPS)*, 1989.
- [11] R. Yu, A. Li, C.-F. Chen, J.-H. Lai, V. I. Morariu, X. Han, M. Gao, C.-Y. Lin, and L. S. Davis, "Nisp: Pruning networks using neuron importance score propagation," 2017.
- [12] B. S and D. UN, "Brain-derived neurotrophic factor and its clinical implications," *Arch Med Sci*. 2015;11(6):1164–1178. doi:10.5114/aoms.2015.56342, 2015.
- [13] N. Shazeer, A. Mirhoseini, K. Maziarz, A. Davis, Q. Le, G. Hinton, and J. Dean, "Outrageously large neural networks: The sparsely-gated mixture-of-experts layer," 2017.
- [14] L. C. J. W. T. Sanh, Victor; Debut, "Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter," *arXiv preprint arXiv:1910.01108*, 2019.
- [15] D. Balduzzi, W. M. Czarnecki, T. W. Anthony, I. M. Gemp, E. Hughes, J. Z. Leibo, G. Piliouras, and T. Graepel, "Smooth markets: A basic mechanism for organizing gradient-based learners," 2020.
- [16] P. Dütting, Z. Feng, H. Narasimhan, D. C. Parkes, and S. S. Ravindranath, "Optimal auctions through deep learning," 2017.

6 Appendix

6.1 Deriving the idealized ranking.

We approximate the change in the benchmark $\Delta \mathcal{B}$ at a local minimum under a perturbation $\Delta F(x)_i = [\dots, -f_i(x), \dots]$ reflecting the removal of the i^{th} node.

$$\Delta \mathcal{B} = \mathcal{B}(F + \Delta F_i) - \mathcal{B}(F) = \sum_i^n \mathcal{L}_i(F + \Delta F_i) - \mathcal{L}_i(F) \quad (17)$$

$$\mathcal{L}_i(F + \Delta F_i) - \mathcal{L}_i(F) \approx \frac{\partial \mathcal{L}_i}{\partial F} * \Delta F_i + \frac{1}{2} \Delta F_i^T * H(\mathcal{L}_i) * \Delta F_i + O(\Delta F_i^3) \quad (18)$$

(17) follows from the definition of the benchmark \mathcal{B} and (18) follows from a Taylor series under the perturbation $\Delta F(x)_i$. Note that the first term $\frac{\partial \mathcal{L}_i}{\partial F}$ is zero at the local minimum and the higher order term $O(\Delta F_i^3)$ can be ignored for sufficiently small perturbations. These assumptions are also made by [10] and [11]. Note that \mathcal{L}_i is an expectation over the dataset D_i , and all terms are evaluated at a point x so we have:

$$\Delta \mathcal{B} \approx \sum_i^n \sum_{x \in D_i} \Delta F_i^T(x) * H(\mathcal{Q}_i(x)) * \Delta F_i(x) \quad (19)$$

Here the hessian over the error function $H(\mathcal{Q}_i(x))$ and the summation over the dataset $\sum_{x \in D_i}$ have been appropriately substituted. The constant factor $\frac{1}{2}$ can be removed and this leaves our result.

6.2 Deriving the weight convergence game.

6.3 Theorem

For choice of Hessians $H(\mathcal{L}(F))$ the network convergence-game can be described with the following linear relationship between gradient terms:

$$\frac{\partial P}{\partial W} = \frac{\alpha}{\tau} * \frac{\partial L}{\partial W} + \frac{\partial R}{\partial W} \quad (20)$$

299

300 With the gradient of the loss:

$$\frac{\partial L}{\partial W} = \frac{\partial}{\partial W} [(F_W - F_{W_0})^T * H(\mathcal{L}(F)) * (F_W - F_{W_0})] \quad (21)$$

301

302 6.4 Setup

303 We analyze the system by characterizing the behaviour of participants via their payoff in two terms:

- 304 1. The utility attached to that participant's loss as a function of their weights $U(L(W))$. U is assumed
- 305 roughly linear for small change in the weight matrix, $U(\mathcal{L}) = \alpha * \mathcal{L}$, and $\frac{\partial U}{\partial \mathcal{L}} = \alpha$.
- 306 2. The network is converged to a local minimum in the inputs $\frac{\partial \mathcal{L}}{\partial F} = 0$.

307 (1) and (2) From the payoff formulation in 6.3 we write:

$$P(W) = \alpha * L(W) + \tau * R(W) \quad (22)$$

308

309 Note, the utility function and emission were measured in similar units and so α is the *price* of each unit change
 310 in loss. The analysis just supposes such a score exists, not that it can be computed. Participants are selecting
 311 their weights by making gradient steps $\Delta W_i = \frac{\partial P_i}{\partial W_i} = \frac{\partial U_i}{\partial W_i} + \tau * \frac{\partial R_i}{\partial W_i}$ as to maximize their local payoff. For
 312 brevity we omit the subscript i for the remainder of the analysis. Consider a Taylor expansion of the loss under a
 313 change ΔF in the inputs.

$$\mathcal{L}(F + \Delta F) = L(F) + \frac{\partial \mathcal{L}}{\partial F} \Delta F + \frac{1}{2} \Delta F * H(\mathcal{L}(F)) * \Delta F + O(\Delta F^3) \quad (23)$$

314 The first linear term $\frac{\partial \mathcal{L}}{\partial F}$ is zero and the higher order terms are removed for sufficiently small perturbations in F .
 315 We then perform a change of variable $F = F_{W_0}$, and $\Delta F = F_{W_1} - F_{W_0}$ where W_0 are the current set of
 316 weights and W_1 are another choice such that F_{W_0} and F_{W_1} are those inputs masked by W_0 and W_1 according
 317 to (7). Substituting this into (23):

$$\mathcal{L}(F_{W_1}) = L(F_{W_0}) + \frac{1}{2} (F_{W_1} - F_{W_0})^T * H(\mathcal{L}(F)) * (F_{W_1} - F_{W_0}) \quad (24)$$

318 The function $\mathcal{L}(F_{W_1})$ is simply an approximation of the loss for any choice of weights W_1 given that the
 319 network has converged under W_0 . Finally, by the α -linear assumption of the utility we can attain the following:

$$\frac{\partial U}{\partial W} = \alpha * \frac{\partial L}{\partial W} = \alpha * \frac{\partial}{\partial W} [(F_W - F_{W_0})^T * H(\mathcal{L}(F)) * (F_W - F_{W_0})] \quad (25)$$

320 Note that we've dropped the subscript W_1 , $L(F_{W_0})$ is constant not depending on the choice of weights, and
 321 the fraction $\frac{1}{2}$ has been removed. The remaining term $\frac{\partial R}{\partial W}$ is derivable via the ranking ranking function in
 322 Section 1.3. Finally, dividing both terms by τ shows the result:

$$\frac{\partial P}{\partial W} \approx \frac{\alpha}{\tau} * \frac{\partial L}{\partial W} + \frac{\partial R}{\partial W} \quad (26)$$

$$\frac{\partial L}{\partial W} = \frac{\partial}{\partial W} [(F_W - F_{W_0})^T * H(\mathcal{L}(F)) * (F_W - F_{W_0})] \quad (27)$$

323

324 6.5 Deriving the ex-post zero-regret step.

325 Consider the system described above. A set of n nodes are changing the weights in the ranking matrix W
 326 iteratively using gradient descent with learning rate λ . $W^{t+1} = W^t + \lambda \Delta W$. Here, the change of weights is
 327 $\Delta W = [\Delta w_0, \dots, \Delta w_n]$ where each Δw_i is the weight change pushed by node i . Each node is attempting to
 328 competitively maximize it's payoff as a function of the weights $P_i(W)$.

329 6.5.1 Definition

330 The ex-post regret for a single step is the maximum difference in loss between the chosen step Δw_i and all
 331 alternative Δw_i^* . The *expected* ex-post regret is this difference in expectation, where the expectation is taken
 332 over all choices Δw_j 's chosen by other participants [16].

$$\text{rgt}_i = E_{\Delta w_j} [\max_{\Delta w_i^*} [P_i(\Delta w_i^*) - P_i(\Delta w_i)]] \quad (28)$$

333 6.5.2 Theorem

334 For sufficiently small λ , the expected ex-post regret for strategy $\Delta w_i = \frac{\partial P}{\partial w_i}$ is 0.

335 6.5.3 Proof

336 Consider Taylor's theorem at the point W for the payoff function P under a change in weights $W^* = W + \lambda \Delta W$.
 337 There exists a function $h(W^*)$ such that in the limit as, $W^* \rightarrow W$ we have the exact equivalence:

$$P(W^*) = P(W) + \frac{\partial P}{\partial W} (W^* - W) + h(W^*) \quad (29)$$

338 Let $P(W_*)$ represent the payoff when the weight change of the i^{th} row is $\Delta W_i = \frac{\partial P}{\partial W_i}$, and let $P(W^*)$ be any
 339 other choice. By the definition of regret, and Taylors theorem as $\lambda \rightarrow 0$, we have:

$$\text{rgt}_i = E_{\Delta W_j} [\max_{\Delta W_i^*} [\frac{\partial P}{\partial W} (W^* - W) - \frac{\partial P}{\partial W} (W_* - W)]] \quad (30)$$

340 This follows by subtracting (29) with choice W^* and W_* . Next, substituting $W^* - W = -\lambda \Delta W$ and expanding
 341 $\frac{\partial P}{\partial W} \Delta W = [\frac{\partial P}{\partial W_0} * \Delta W_0, \dots, \frac{\partial P}{\partial W_n} * \Delta W_n]$ into the equation above leaves:

$$\begin{aligned} \frac{\partial P}{\partial W} (W^* - W) - \frac{\partial P}{\partial W} (W_* - W) &= \lambda (\frac{\partial P}{\partial W_i} * \Delta W_i^* - \frac{\partial P}{\partial W_i} * \Delta W_{i*}) + \\ &\quad \lambda \sum_{j \neq i}^n (\frac{\partial P}{\partial W_j} * \Delta W_j^* + \frac{\partial P}{\partial W_j} * \Delta W_{j*}) \end{aligned} \quad (31)$$

342 The constant λ can be removed and the second term depends only on weights of other rows $W_{j \neq i}$. These can be
 343 removed under the expectation $E_{\Delta W_j}$. He have:

$$\text{rgt}_i = E_{\Delta W_j} [\max_{\Delta W_i^*} [\frac{\partial P}{\partial W_i} * \Delta W_i^* - \frac{\partial P}{\partial W_i} * \Delta W_{i*}]] \quad (32)$$

344 Finally, we use the the fact that for vectors a, b and angle between them θ the magnitude of the dot product is
 345 $|a||b|\cos\theta$. This is maximized when the vectors are parallel $\theta = 0$ and $\cos(\theta) = 1$, or $\Delta W_i = \kappa * \frac{\partial P}{\partial W_i}$ for
 346 some constant $\kappa > 0$. Thus $P(\Delta W^*)$ is maximize when $\Delta W_i^* = \kappa * \frac{\partial P}{\partial W_i}$. Since $P(\Delta W^*) = P(\Delta W_*)$ in
 347 the maximum, this proves the point.