

---

# BitTensor: A Peer-to-Peer Intelligence Benchmark

---

Anonymous Author(s)  
Affiliation  
Address  
email

## Abstract

1       The dominant tools used to guide machine intelligence are benchmarks: datasets  
2       of labelled examples which rank systems on a set of predefined tasks. We argue,  
3       however, that task based benchmarks are inefficient guides for machine intelligence.  
4       Instead, we propose a benchmark which ranks *knowledge production* from within  
5       a culture of intelligence systems. In this paper we describe how this benchmark  
6       is constructed and how it is negotiated by computers that share knowledge peer-  
7       to-peer (P2P) across the internet. Since this internet based negotiation introduces  
8       game-theoretic considerations, we design methods to ensure the veracity of the  
9       benchmark when peers remain self-interested and trustless. We then test our design  
10      by empirically showing that the negotiation closely matches the desired ranking at  
11      the system’s competitive equilibrium. The result is a benchmark which is directly  
12      aligned with knowledge production. It is furthermore, inherently collaborative,  
13      open, decentralized, and can consume resources connected by an internet fabric.

14     Increasingly, state-of-the-art machine learning systems focus on the production of *machine knowledge*,  
15     non-task specific understanding of inputs, which can be later tuned to a wide variety of problems [2].  
16     This shift is a consequence of the fact that a focus on tasks during training produces narrow specialists  
17     rather than resilient generalists [1]. In contrast, while becoming the dominant technique for training  
18     individual models, the field still opts for task based objectives as the method of reward at the highest  
19     level of its evolution [3]. The misalignment between this salient feature (machine knowledge) and its  
20     feedback mechanism (task performance) is inefficient for the following reasons:

- 21       1. A definition of intelligence provided by tasks performance produces a feedback mechanism  
22       which converges the field towards this narrow task definition. Unfortunately, simply ex-  
23       panding the set tasks to cover the general problem would require a near infinite number of  
24       supervised (dataset, task) pairs [1]. Because of this, the feedback mechanism at the highest  
25       level of the field remains misaligned with general intelligence [4].
- 26       2. Representational knowledge can be measured in a much higher dimensional space than the  
27       low-dimensional scores produced by the tasks. For instance, the 1024 dimensional floating  
28       point embedding produced by BERT in comparison to the 8 dimensional score produced by  
29       SuperGLUE. In the low dimensional space, we cannot easily apply divide-and-conquer to  
30       the problem to reward niche machine learning systems for working on specific aspects of  
31       understanding. This is wasteful as those models may have been costly to create in the first  
32       place.
- 33       3. Because we do not measure how much one team’s model improved the collective. Task  
34       based benchmarks produce a winner-take-all competition which prunes all but the best. The  
35       result is a world-wide challenge where only those with the means and the resources to train  
36       the largest models may contribute to the state-of-the-art.

37 A benchmark that directly measures machine intelligence would be very beneficial.<sup>1</sup> By analogy,  
 38 similar to rewarding domain specialists or generalists in human cultures, not against their ability  
 39 to solve tasks, but their ability to improve overall understanding. Our proposal is a framework in  
 40 which models directly share knowledge that they have learned with each other. Their peers select that  
 41 knowledge for its ability to improve their own understanding. Crucially, each model’s performance is  
 42 measured by its ability to improve collective understanding of inputs rather than against any specific  
 43 task. The higher resolution of measurement allows the system to rewarded models for improving  
 44 small regions of the general problem. Inherently the benchmark introduces knowledge sharing and  
 45 the demand for performance guides a de-facto ‘market for intelligence’.

46 Practically, we design the benchmark to run in a continuous and asynchronous fashion, peer-to-peer  
 47 (P2P) across the wide network. Since this introduces trustless computation<sup>2</sup> we dedicate a part of  
 48 this paper to explaining how the system remains fair when little assurance can be given about the  
 49 computers that compose it. The contributions of this paper are thus four-fold:

- 50 1. We introduce a P2P based intelligence framework.
- 51 2. We introduce an “intelligence market” that rewards models that improve representational  
 52 knowledge within it.
- 53 3. We show how models can share knowledge to increase each other’s performance which  
 54 removes the need for others to re-learn that knowledge.
- 55 4. We show how self-interested desire to minimize a loss function can be used to guide the  
 56 network towards a fair ranking.

57 The rest of this paper is organized as follows: Section 1 presents the intelligence benchmark and  
 58 ranking mechanism that models use to evaluate each other. Section 2 considers the scenario where  
 59 participants are not honestly reporting the significance of their peers, and proposes a method to  
 60 detect such behavior. Section 3 highlights the experiments performed to verify the function of the  
 61 framework. Finally, Section 4 presents a summary and future works on this system.

## 62 1 Methodology

### 63 1.1 Benchmark

64 The network is composed by  $n$  unique parameterized functions  $F = f_0, \dots, f_j, \dots, f_n$  where each  
 65 function is producing an output tensor  $f_i(F(x))$ , a “representation” from an input tensor  $F(x) =$   
 66  $[f_0(x) \dots f_n(x)]$  gathered by querying its neighbors. Each function is training asynchronously over a  
 67 dataset  $D_i = [X, Y]$  such that, given an error function  $\mathcal{Q}_i$ , its expectation over that data  $E_{D_i}$  defines  
 68 a loss  $\mathcal{L}_i = E_{D_i}[\mathcal{Q}_i(y, f_i(F(x)))]$ . We assume these losses are measured on the same scale and  
 69 thus our benchmark  $\mathcal{B}$  can be defined by their sum:

$$\mathcal{B} = \sum_i^n \mathcal{L}_i \quad (1)$$




Figure 1:  $n = 6$  parameterized functions with losses  $\mathcal{L}_i$  and datasets  $D_i$ .

<sup>1</sup>“The iron rule of nature is: you get what you reward for. If you want ants to come, you put sugar on the floor.” - Charlie Munger

<sup>2</sup>In decentralized systems, trust is shifted from the individuals to the network itself, “trustless” usually means “minimal trust is required” with respect to the protocol

Each parameterized function is represented here in its most abstract sense[6] and need only accept the same input type  $x$  and produce the same output dimension to fit within the network. For instance, unicode encoded string  $x = \text{"hello"}$  and its semantic representation as a  $64 \times 49$  word embedding. This widened scope ensures participants can be multi-task [7], use completely distinct computing substrates [8] or train on unique datasets. [9].

## 1.2 Ideal Ranking

Our goal in this work is to produce a ranking  $R = [R_i]$  over these functions where the score  $R_i \in R$  represents participant  $i$ 's information-theoretic significance to the benchmark  $\mathcal{B}$ . Following Le Cun and others [10, 11], it is reasonable to analytically define this significance by equating it with the cost of removing each component from the network:

$$R_i \approx \frac{1}{n} \sum_j \sum_{x \in D_j} \Delta F^T(x)_i * H(\mathcal{Q}_j(x)) * \Delta F(x)_i \quad (2)$$

$$\Delta F(x)_i = [0, \dots, 0, -f_i(x), 0, \dots, 0]$$

Where the above is derived using a Taylor series (Appendix 6.1) and  $\Delta F(x)_i$  is the perturbation of the  $i^{\text{th}}$  node's inputs as a function of it's choice of weights when removing function  $f_i$  at the point  $x$ . Note, the linear and higher order terms of the Taylor series have been removed following [11] and the remaining term  $H(\mathcal{Q}_i)$  is the hessian of our error function. When the error function  $\mathcal{Q}$  is the twice-differentiable cross-entropy, then  $H(\mathcal{Q}_i)$  is the Fisher information matrix, and  $R_i \in R$  is measured as relative entropy: reflects each participants informational significance to the network as a whole.

## 1.3 Inter Ranking

It is not possible to compute the ranking score above without access to the parameters of each function in the network. Instead, we use a set of inter-model weights  $W = [w_{ij}]$  where each  $w_{i,j}$  is the score attributed to  $f_j$  from  $f_i$  combined into an  $n \times n$  square matrix.

$$w_{ij} = \sum_{x \in D_i} \Delta F^T(x)_j * H(\mathcal{Q}_i(x)) * \Delta F(x)_j \quad (3)$$




Figure 2: Inter-model contribution weights:  $w_{i,j}$  the score attributed to  $f_j$  from  $f_i$

The weights can be computed on the fly either approximately [11] or by using the full hessian of the error. We store them on a distributed ledger and allow participants to update them by making changes of bounded size:  $W^{t+1} = W^t + \lambda \Delta W$ , where  $\|W_i\|_2 < \epsilon$  at block step  $t$ . We also enforce that the scores in each row sum to 1,  $\|w\|_1 = 1$ . The reasoning for this enforcement is to ensure stake fairness between early movers and newcomers with innovations that enable them to perform better overall.

97 The equivalent ranking  $R$  (2) can then be computed by normalizing column sum of the weight matrix:  
 98

$$R = \frac{1}{n} W^T * \mathbb{1} \quad (4)$$

99 The problem is that without system-wide access to the model parameters, the computation of  $w_{ij}$  (3)  
 100 is intractable, as it would require collecting great amounts of information from potentially billions of  
 101 neurons. It is reasonable to assume participants will select weights which artificially increase their  
 102 own rank rather than others in the network. Moreover, since the network remains open, participants  
 103 may choose to create many spuriously neighbours and rank themselves higher. The remainder of this  
 104 paper describes our proposal for resolving these issues.

## 105 1.4 Stake

106 The proposed solution begins by introducing a finite resource  $S = [s_i]$ , a component's 'stake' in the  
 107 system, and an inflation mechanism  $\tau$  which translates the ranking vector  $R$  into additional stake as  
 108 incentive.

$$R = \frac{1}{n} W^T \circ S * \mathbb{1} \quad (5)$$

$$S^{t+1} = S^t + \tau * \frac{R}{||R||_2} \quad (6)$$

109 The  $\circ$  here is the Hadamard product between the  $n \times n$  matrix  $W$  and the  $n \times n$  matrix containing  $S$   
 110 in each column, and  $t$  is the time-step referred to in Section 1.3 (measured in distinct blocks on the  
 111 distributed ledger). By design (5) increases the importance of those with stake,  $s_0 * w_{ij}$ . This serves  
 112 two purposes:

- 113 1. New computers can spuriously create new nodes, however this won't game the ranking  
 114 because the amount of stake they hold is finite.
- 115 2. The resource provides mechanism power.

116 By providing it to nodes with large rank, this ensures that those with weight must have worked to  
 117 attain it, or indirectly subsidized those who have done so already. A single staked token would be  
 118 enough to bootstrap the process. <sup>3</sup>

---

### Algorithm 1 Inflation mechanism

---

**Require:**  $S = [n \times 1]$  Stake Vector  
**Require:**  $W = [n \times n]$  Weight Matrix  
**Require:**  $\tau > 0$  inflation rate  
**while** TRUE **do**  
    $W = W + \lambda \Delta W$   
    $R = \frac{1}{n} W^T \circ S * \mathbb{1}$   
    $S = S + \tau * \frac{R}{||R||_2}$   
**end while**

---

<sup>3</sup>The propagation of a resource mocks the flow of the brain-derived neurotrophic factor (BDNF). [12]

## 1.5 Competitive weights

While stake provides some protection against malicious actors, it does not ensure weights are set accurately. Our solution begins by introducing competition for connectivity within the network. Nodes that underweight are punished by having inputs from the network masked to zero (7), **which in turn would punish the loss function**. To frame this market we borrow the continuous differential activation function  $\sigma$  with range  $(0, 1)$ . Under a choice of weights  $W_i$  the inputs to component  $i$  are:

$$F_W(x) = [f_0(x) * \sigma(s_i * w_{i,0} - \mu_0), \dots, f_n(x) * \sigma(s_i * w_{i,n} - \mu_n)] \quad (7)$$

$$\sigma = \frac{1}{1 + e^{-\frac{x}{T}}} \quad (8)$$

Here, the shift term  $\mu_j$  is the average of the weights in each column  $\mu_j = (\frac{1}{n}) \sum_i s_i * w_{i,j}$ , and the activation function is the temperature scaled sigmoid. Because the allocation mechanism is standard across the network it is possible for each participants to compute both  $\frac{\partial \mathcal{L}_i}{\partial W_i}$  and  $\frac{\partial R_i}{\partial W_i}$ . Computers may augment their usual training framework, for instance, Tensorflow, with the allocation mechanism shown here.

## 1.6 Running the network

The steps to run a network participant are:

1. Participant defines its dataset  $D_i$ , loss  $\mathcal{L}_i$  and parameterized function  $f_i$
2. At each training iteration the participant broadcasts batches of examples from  $D_i$  to its peers  $[batch\_size, x]$
3. Responses  $F(x)$  from the network produce a loss-gradient  $\frac{\partial \mathcal{L}}{\partial F}$  which back-propagates through  $f_i$  and out to the network.
4. During 2 and 3 the participant competitively selects the weights for their row  $w_{ij} \in W$ .
5. Participants submit changes to the weights  $\Delta W_i$  which, in-turn changes the ranking and induces inflation  $\tau * R$ .
6. **After steps 1-5 have been performed multiple times**, participants disconnect and verify the model in a normal manner.

Peers only communicate with computers that hold stake as a consequence of section 1.5. Those that fail to produce value will be pruned naturally as participants learn to differentiate signal from noise.

## 1.7 Conditional computation

As the network grows, outward bandwidth will become the major bottleneck. Components learn to trim outward bandwidth by employing a Sparsely-Gated Mixture-of-Experts (SGMoE) [13] layer at the input. The gating layer determines a sparse combination of children to query for each example and then re-joins them using the the gating weights  $g_j(x)$ . The combined gated inputs are fed as input to the local function:

$$f_i = f_i(G(x)) \quad (9)$$

$$G(x) = [\dots, g_j(x) * f_j(x), \dots] \quad (10)$$

The layer cuts outward bandwidth, querying only a small subset of peers for each example. The gating function is trainable w.r.t to the loss and its weights act as a proxy for importance  $w_{ij} \in W$ . This method has been shown to drastically increase the potential for outward bandwidth in datacenter training,[13] and has been investigated in a peer-to-peer (P2P) setting as well [5]

## 1.8 Extracting knowledge

Inter-node dependence in the network is broken using distillation[6], a compression and knowledge technique in which a smaller model – the student – mimics the behaviour of an ensemble. We employ this technique over the gating ensemble (10) where the student model learns to minimize the cross-entropy (shown below as KL) between the logits produced by the gating network and its predicted distribution. [14]

$$\text{distillation loss} = \text{KL}_D(\text{dist}(x), G(x)) \quad (11)$$

We use the distilled model as proxy to cut recursive calling between each components rather than query farther into the network. If models go offline, their peers can use their distilled versions in-place. Private data can be validated over the distilled models instead of querying the network. Eventually, components can fully disconnect from the network using the distilled inputs to validate and inference the models offline.




Figure 3: Queries propagate to depth=1 before the distilled model is used.

## 2 Analysis

We consider the scenario where participants are not honestly reporting the significance of their peers. The network is progressing at discrete timesteps  $t$  and it is reasonable to assume that each participant is attempting to maximize their subjective payoff over these steps.

### 2.1 Payoff Model

The staking system in Section 1.4 gives participants incentive to maximize their self-weight  $w_{ii}$  while the competitive connectivity described in Section 1.5 makes it costly for participants to decrease the remaining weights  $w_{ij}$  in their row. Since the row must sum to 1, we have a trade-off in two terms:

1. A utility term attached to the loss  $U(\mathcal{L}(W))$ .
2. The token emission via inflation  $\tau * R_i(W)$ .

Both of these are functions of the weights and, without loss of generality, measured in similar units:

$$P_i(W) = U_i(\mathcal{L}_i(W)) + \tau * R(W)_i \quad (12)$$

It is reasonable to assume payoff maximizing participants will use  $P_i$  as their objective during training. This can be computed using standard tools since both terms  $U(\mathcal{L}_i(W))$  and  $R(W)_i$  are fully continuous and differentiable. The system can be characterized as a competitive gradient descent game where participants are making steps  $\Delta W_i = \frac{\partial P_i}{\partial W_i}$ . In appendix 6.2 we prove that this strategy

181 is regret-free and achieves the best expected payoff in hindsight. This assumption is also generally  
 182 employed in smooth markets [15]. The iterative descent thus follows:

$$W^{t+1} = W^t + \lambda \Delta W \quad (13)$$

183

$$\Delta W = [\frac{\partial P_0}{\partial W_0}; \dots; \frac{\partial P_n}{\partial W_n}] \quad (14)$$

184

## 185 2.2 Empirical Model

186 Without access to the running network we evaluate our system using an empirical model. To derive  
 187 the gradient steps in this model  $\frac{\partial P_i}{\partial W_i}$  we make the following assumptions:

- 188 1. The utility functions are continuous-differentiable and are bounded by their first order  
 189 derivatives  $\frac{\partial U}{\partial \mathcal{L}} = \alpha$ .
- 190 2. The network is converged to a local minimum in the inputs  $\frac{\partial \mathcal{L}}{\partial F} = 0$ .

191 The first assumption is approximate for small changes of continuous functions and the second  
 192 assumption is realistic after extended training. In Appendix 6.2 we derive the following gradient step:  
 193

$$\frac{\partial P}{\partial W} = \frac{\alpha}{\tau} * \frac{\partial L}{\partial W} + \frac{\partial R}{\partial W} \quad (15)$$

194

$$\frac{\partial L}{\partial W} = \frac{\partial}{\partial W} [(F_W - F_{W_0})^T * H(\mathcal{L}(F)) * (F_W - F_{W_0})] \quad (16)$$

195

196 Here  $F_W$  is the masked inputs from Section 1.5.  $(F_W - F_{W_0})$  is the difference in the mask between  
 197 the choice of weights  $W$  and the weights at the minimum  $W_0$  and  $H(\mathcal{L}(F))$  is the  $[n \times n]$  hessian of  
 198 the loss over inputs  $F$ . This formulation is both intuitive and useful: the gradient term  $\frac{\partial L}{\partial W}$  measures  
 199 the change in loss for any choice of weights. Additionally, it also allows us to compute the ranking  
 200 for simulated hessian terms.

201 The remainder of the network is deterministic, and can be described by a choice of  $\theta =$   
 202  $[\alpha, \tau, \lambda, n, \sigma, T, S, W, H]$ . For instance, the secondary term,  $\frac{\partial R}{\partial W}$  can be computed-directly from  
 203 Section 1.4 and only depends on the stake vector  $S$  and weights  $W$ .

## 204 3 Experiments

205 To generate sample statistics from the network, we first select  $[\alpha, \tau, \lambda, S, \sigma, T, n] \in \theta$ , then we  
 206 generate random positive semi-definite Hessians  $H$ , and random uniform initial weights  $W_0$ . For each  
 207 parameterization we discover the competitive ranking by converging the system to a Nash-equilibrium:  
 208 an equilibrium where no individual can vary from their set of weights and stand to gain. [16] To find  
 209 this equilibrium, we use the competitive descent strategy described in (14) and compute the gradient  
 210 terms from (15). In each trial we use a learning rate  $\lambda = 0.005$  and stop when the gradient terms are  
 211 bounded by  $\epsilon$  or the steps exceed  $1000 \times n$ . The competitive ranking  $R^*$  at this point follows from  
 212 (4) and can be compared to the idealized score in  $R$  from (2).

213 We are interested in the correlation between  $R^*$  and  $R$  as we vary the ratio  $\alpha$  and  $\tau$ , this ratio is  
 214 explicit in (15) where the fraction relates the two gradient terms. Intuitively the ratio is between the

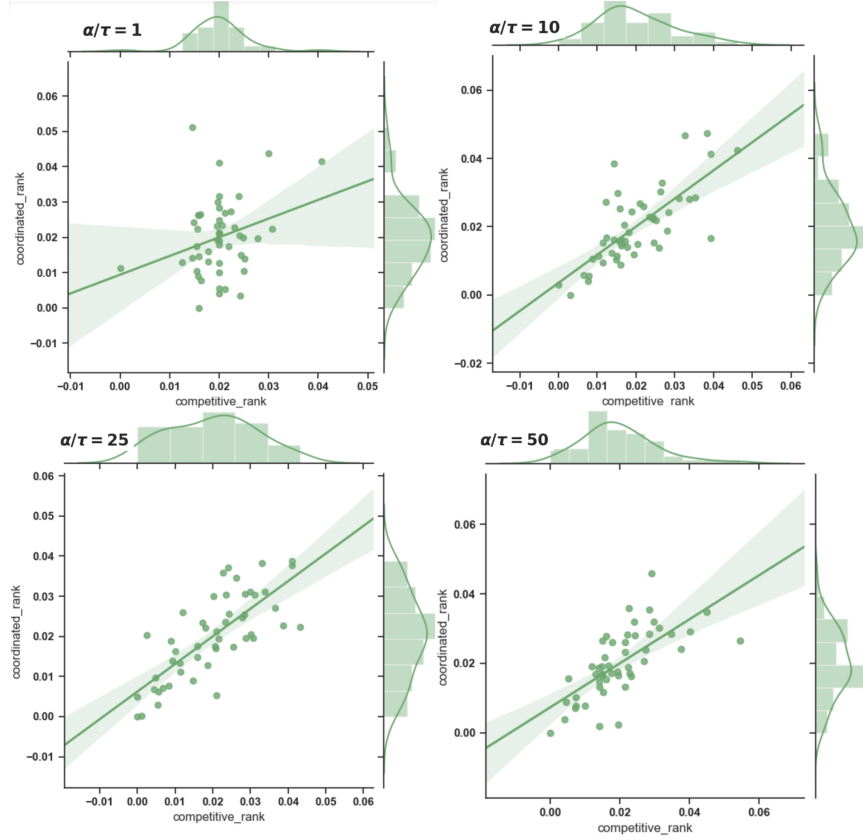


Figure 4: Correlations between the competitive rank and coordinated rank for  $\frac{\alpha}{\tau} \in \{1, 10, 25, 50\}$ . For low values of  $\frac{\alpha}{\tau}$  the weights converge to the identity: the state where peers are fully disconnected.




Figure 5: (i) Left: the ratio between the main diagonal and the remainder of the weights. Right: the adaptive  $\tau$  parameter converging onto the target. The weight matrix sparsity is a proxy for the ranking accuracy which we see in Figure 4. As sparsity converges onto the target the ranking is also improving.

215 value of minimizing the loss and maximizing revenue from inflation. Since this effects the ranking,  
 216 we show this trade-off for various choices in Figure 4.



Finally, in Figure-5 we implement an adaptive- $\tau$  strategy where the network varies the inflation rate. Initial inflation is zero and then increases until the weights begins to converge towards the main diagonal  $w_{ii} = 1$ . We measure the sparsity  $sparsity = \text{sum}(W_{dg})/\text{sum}(W)$ , the ratio between the main diagonal and the remaining weights. As sparsity increases we push the market equilibrium by decreasing  $\tau$ . Figure shows this adaptive convergence for  $\alpha = 1$  with a sparsity target of 1.

### 3.1 Discussion

Figure 4 shows the relationship between the idealized rank and the competitive rank as a function of  $\frac{\alpha}{\tau}$ . Figure-4-a,  $\frac{\alpha}{\tau} = 1$  shows the case where this ratio negatively effects the ranking accuracy. Here, all components have set  $w_{ii} = 1$  and the resulting scores for all participants has converged to  $1/n$ . When this occurs, the system could decrease the inflation rate  $\tau$  and push the network towards the high information markets seen in Figure-4-b, c, and d. Figure-5 shows a basic implementation of this where  $\tau$  adapts to the ratio between the row-sums and the main-diagonal. By lowering inflation, it subsequently 'costs less' to connect with peers. Profit maximizing nodes automatically adjust to the change and the system converges back towards an accurate ranking. Those with high ranks will oppose inflation decreased, while those with low ranks will welcome it. The equilibrium found in this meta game will most certainly depend on the number of participants, key to both the ranking accuracy and the market at its core. However, we leave this analysis for a follow up paper.

## 4 Conclusion

We have proposed an inter-model benchmark that can run in a P2P setting outside of a trusted environment. We started with a typical machine learning framework defined by a set of functions with their losses over given datasets, then derived an idealized ranking score. The measure produced an information theoretic score that new participants can improve by learning how to be useful to their peers. However, the system is incomplete without a mechanism that prevents participants from ranking dishonestly. To resolve this, we proposed an incentive scheme and a differential allocation system over the network weights. The system allows the participants to train for connectivity in the graph. Following this, we described how to increase the outward bandwidth in the system using a trainable gating network and how to cut dependence between nodes using distillation. Finally we showed how increasing the number of nodes in the system and fixing the inflation mechanism properly ensured that the resulting rank scores would correlate with those found in a idealized setting. While this is true, stake in the system holds value as a means to drive what the network learns. That benchmark is continually being solved by the participants, compounding what has been learned before and making it available to new learners in the system.

## 5 Broader Impact

In order to provide a balanced perspective, authors are required to include a statement of the potential broader impact of their work, including its ethical aspects and future societal consequences. Authors should take care to discuss both positive and negative outcomes.

## References

- [1] A. Radford, J. Wu, R. Child, D. Luan, D. Amodei, and I. Sutskever, "Language models are unsupervised multitask learners," *OpenAI Blog*, vol. 1, no. 8, p. 9, 2019.
- [2] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, "Bert: Pre-training of deep bidirectional transformers for language understanding," *arXiv preprint arXiv:1810.04805*, 2018.
- [3] A. Wang, A. Singh, J. Michael, F. Hill, O. Levy, and S. R. Bowman, "Glue: A multi-task benchmark and analysis platform for natural language understanding," *arXiv preprint arXiv:1804.07461*, 2018.
- [4] F. Chollet, "On the measure of intelligence," *arXiv preprint arXiv:1911.01547*, 2019.
- [5] M. Riabinin and A. Gusev, "Learning@home: Crowdsourced training of large neural networks using decentralized mixture-of-experts," *arXiv preprint arXiv:2002.04013*, 2020.
- [6] G. Hinton, O. Vinyals, and J. Dean, "Distilling the knowledge in a neural network," *arXiv preprint arXiv:1503.02531*, 2015.

- [7] L. Kaiser, A. N. Gomez, N. Shazeer, A. Vaswani, N. Parmar, L. Jones, and J. Uszkoreit, “One model to learn them all,” 2017.
- [8] M. A. Nugent and T. W. Molter, “Cortical processing with thermodynamic-ram,” 2014.
- [9] G. Lample and A. Conneau, “Cross-lingual language model pretraining,” 2019.
- [10] Y. LeCun, D. J. S., and S. S. A., “Optimal brain damage,” *Advances in Neural Information Processing Systems 2 (NIPS)*, 1989.
- [11] R. Yu, A. Li, C.-F. Chen, J.-H. Lai, V. I. Morariu, X. Han, M. Gao, C.-Y. Lin, and L. S. Davis, “Nisp: Pruning networks using neuron importance score propagation,” 2017.
- [12] B. S and D. UN, “Brain-derived neurotrophic factor and its clinical implications,” *Arch Med Sci*. 2015;11(6):1164–1178. doi:10.5114/aoms.2015.56342, 2015.
- [13] N. Shazeer, A. Mirhoseini, K. Maziarz, A. Davis, Q. Le, G. Hinton, and J. Dean, “Outrageously large neural networks: The sparsely-gated mixture-of-experts layer,” 2017.
- [14] L. C. J. W. T. Sanh, Victor; Debut, “Distilbert, a distilled version of bert: smaller, faster, cheaper and lighter,” *arXiv preprint arXiv:1910.01108*, 2019.
- [15] D. Balduzzi, W. M. Czarnecki, T. W. Anthony, I. M. Gemp, E. Hughes, J. Z. Leibo, G. Piliouras, and T. Graepel, “Smooth markets: A basic mechanism for organizing gradient-based learners,” 2020.
- [16] P. Dütting, Z. Feng, H. Narasimhan, D. C. Parkes, and S. S. Ravindranath, “Optimal auctions through deep learning,” 2017.

## 6 Appendix

### 6.1 Deriving the idealized ranking.

We approximate the change in the benchmark  $\Delta\mathcal{B}$  at a local minimum under a perturbation  $\Delta F(x)_i = [\dots, -f_i(x), \dots]$  reflecting the removal of the  $i^{th}$  node.

$$\Delta\mathcal{B} = \mathcal{B}(F + \Delta F_i) - \mathcal{B}(F) = \sum_i^n \mathcal{L}_i(F + \Delta F_i) - \mathcal{L}_i(F) \quad (17)$$

$$\mathcal{L}_i(F + \Delta F_i) - \mathcal{L}_i(F) \approx \frac{\partial \mathcal{L}_i}{\partial F} * \Delta F_i + \frac{1}{2} \Delta F_i^T * H(\mathcal{L}_i) * \Delta F_i + O(\Delta F_i^3) \quad (18)$$

(17) follows from the definition of the benchmark  $\mathcal{B}$  and (18) follows from a Taylor series under the perturbation  $\Delta F(x)_i$ . Note that the first term  $\frac{\partial \mathcal{L}_i}{\partial F}$  is zero at the local minimum and the higher order term  $O(\Delta F_i^3)$  can be ignored for sufficiently small perturbations. These assumptions are also made by [10] and [11]. Note that  $\mathcal{L}_i$  is an expectation over the dataset  $D_i$ , and all terms are evaluated at a point  $x$  so we have:

$$\Delta\mathcal{B} \approx \sum_i^n \sum_{x \in D_i} \Delta F_i^T(x) * H(\mathcal{Q}_i(x)) * \Delta F_i(x) \quad (19)$$

Here the hessian over the error function  $H(\mathcal{Q}_i(x))$  and the summation over the dataset  $\sum_{x \in D_i}$  have been appropriately substituted. The constant factor  $\frac{1}{2}$  can be removed and this leaves our result.

### 6.2 Deriving the weight convergence game.

#### 6.3 Theorem

For choice of Hessians  $H(\mathcal{L}(F))$  the network convergence-game can be described with the following linear relationship between gradient terms:

$$\frac{\partial P}{\partial W} = \frac{\alpha}{\tau} * \frac{\partial L}{\partial W} + \frac{\partial R}{\partial W} \quad (20)$$

With the gradient of the loss:

$$\frac{\partial L}{\partial W} = \frac{\partial}{\partial W} [(F_W - F_{W_0})^T * H(\mathcal{L}(F)) * (F_W - F_{W_0})] \quad (21)$$

## 6.4 Setup

We analyze the system by characterizing the behaviour of participants via their payoff in two terms:

1. The utility attached to that participant's loss as a function of their weights  $U(L(W))$ .  $U$  is assumed roughly linear for small change in the weight matrix,  $U(\mathcal{L}) = \alpha * \mathcal{L}$ , and  $\frac{\partial U}{\partial \mathcal{L}} = \alpha$ .
2. The network is converged to a local minimum in the inputs  $\frac{\partial \mathcal{L}}{\partial F} = 0$ .

(1) and (2) From the payoff formulation in 6.3 we write:

$$P(W) = \alpha * L(W) + \tau * R(W) \quad (22)$$

Note, the utility function and emission were measured in similar units and so  $\alpha$  is the *price* of each unit change in loss. The analysis just supposes such a score exists, not that it can be computed. Participants are selecting their weights by making gradient steps  $\Delta W_i = \frac{\partial P_i}{\partial W_i} = \frac{\partial U_i}{\partial W_i} + \tau * \frac{\partial R_i}{\partial W_i}$  as to maximize their local payoff. For brevity we omit the subscript  $i$  for the remainder of the analysis. Consider a Taylor expansion of the loss under a change  $\Delta F$  in the inputs.

$$\mathcal{L}(F + \Delta F) = L(F) + \frac{\partial \mathcal{L}}{\partial F} \Delta F + \frac{1}{2} \Delta F * H(\mathcal{L}(F)) * \Delta F + O(\Delta F^3) \quad (23)$$

The first linear term  $\frac{\partial \mathcal{L}}{\partial F}$  is zero and the higher order terms are removed for sufficiently small perturbations in  $F$ . We then perform a change of variable  $F = F_{W_0}$ , and  $\Delta F = F_{W_1} - F_{W_0}$  where  $W_0$  are the current set of weights and  $W_1$  are another choice such that  $F_{W_0}$  and  $F_{W_1}$  are those inputs masked by  $W_0$  and  $W_1$  according to (7). Substituting this into (23):

$$\mathcal{L}(F_{W_1}) = L(F_{W_0}) + \frac{1}{2} (F_{W_1} - F_{W_0})^T * H(\mathcal{L}(F)) * (F_{W_1} - F_{W_0}) \quad (24)$$

The function  $\mathcal{L}(F_{W_1})$  is simply an approximation of the loss for any choice of weights  $W_1$  given that the network has converged under  $W_0$ . Finally, by the  $\alpha$ -linear assumption of the utility we can attain the following:

$$\frac{\partial U}{\partial W} = \alpha * \frac{\partial L}{\partial W} = \alpha \frac{\partial}{\partial W} [(F_W - F_{W_0})^T * H(\mathcal{L}(F)) * (F_W - F_{W_0})] \quad (25)$$

Note that we've dropped the subscript  $W_1$ ,  $L(F_{W_0})$  is constant not depending on the choice of weights, and the fraction  $\frac{1}{2}$  has been removed. The remaining term  $\frac{\partial R}{\partial W}$  is derivable via the ranking ranking function in Section 1.3. Finally, dividing both terms by  $\tau$  shows the result:

$$\frac{\partial P}{\partial W} \approx \frac{\alpha}{\tau} * \frac{\partial L}{\partial W} + \frac{\partial R}{\partial W} \quad (26)$$

$$\frac{\partial L}{\partial W} = \frac{\partial}{\partial W} [(F_W - F_{W_0})^T * H(\mathcal{L}(F)) * (F_W - F_{W_0})] \quad (27)$$

## 6.5 Deriving the ex-post zero-regret step.

Consider the system described above. A set of  $n$  nodes are changing the weights in the ranking matrix  $W$  iteratively using gradient descent with learning rate  $\lambda$ .  $W^{t+1} = W^t + \lambda \Delta W$ . Here, the change of weights is  $\Delta W = [\Delta w_0, \dots, \Delta w_n]$  where each  $\Delta w_i$  is the weight change pushed by node  $i$ . Each node is attempting to competitively maximize its payoff as a function of the weights  $P_i(W)$ .

### 6.5.1 Definition

The ex-post regret for a single step is the maximum difference in loss between the chosen step  $\Delta w_i$  and all alternative  $\Delta w_i^*$ . The *expected* ex-post regret is this difference in expectation, where the expectation is taken over all choices  $\Delta w_j$ 's chosen by other participants [16].

$$\text{rgt}_i = E_{\Delta w_j} [\max_{\Delta w_i^*} [P_i(\Delta w_i^*) - P_i(\Delta w_i)]] \quad (28)$$

### 331 6.5.2 Theorem

332 For sufficiently small  $\lambda$ , the expected ex-post regret for strategy  $\Delta w_i = \frac{\partial P}{\partial w_i}$  is 0.

### 333 6.5.3 Proof

334 Consider Taylor's theorem at the point  $W$  for the payoff function  $P$  under a change in weights  $W^* = W + \lambda \Delta W$ .

335 There exists a function  $h(W^*)$  such that in the limit as,  $W^* \rightarrow W$  we have the exact equivalence:

$$P(W^*) = P(W) + \frac{\partial P}{\partial W}(W^* - W) + h(W^*) \quad (29)$$

336 Let  $P(W_*)$  represent the payoff when the weight change of the  $i^{th}$  row is  $\Delta W_i = \frac{\partial P}{\partial W_i}$ , and let  $P(W^*)$  be any  
337 other choice. By the definition of regret, and Taylors theorem as  $\lambda \rightarrow 0$ , we have:

$$\text{rgt}_i = E_{\Delta W_j} [\max_{\Delta W_i^*} [\frac{\partial P}{\partial W}(W^* - W) - \frac{\partial P}{\partial W}(W_* - W)]] \quad (30)$$

338 This follows by subtracting (29) with choice  $W^*$  and  $W_*$ . Next, substituting  $W^* - W = -\lambda \Delta W$  and expanding

339  $\frac{\partial P}{\partial W} \Delta W = [\frac{\partial P}{\partial W_0} * \Delta W_0, \dots, \frac{\partial P}{\partial W_n} * \Delta W_n]$  into the equation above leaves:

$$\begin{aligned} \frac{\partial P}{\partial W}(W^* - W) - \frac{\partial P}{\partial W}(W_* - W) &= \lambda (\frac{\partial P}{\partial W_i} * \Delta W_i^* - \frac{\partial P}{\partial W_i} * \Delta W_{i*}) + \\ &\quad \lambda \sum_{j \neq i}^n (\frac{\partial P}{\partial W_j} * \Delta W_j^* + \frac{\partial P}{\partial W_j} * \Delta W_{j*}) \end{aligned} \quad (31)$$

340 The constant  $\lambda$  can be removed and the second term depends only on weights of other rows  $W_{j \neq i}$ . These can be  
341 removed under the expectation  $E_{\Delta W_j}$ . He have:

$$\text{rgt}_i = E_{\Delta W_j} [\max_{\Delta W_i^*} [\frac{\partial P}{\partial W_i} * \Delta W_i^* - \frac{\partial P}{\partial W_i} * \Delta W_{i*}]] \quad (32)$$

342 Finally, we use the the fact that for vectors  $a, b$  and angle between them  $\theta$  the magnitude of the dot product is  
343  $|a||b|\cos\theta$ . This is maximized when the vectors are parallel  $\theta = 0$  and  $\cos(\theta) = 1$ , or  $\Delta W_i = \kappa * \frac{\partial P}{\partial W_i}$  for  
344 some constant  $\kappa > 0$ . Thus  $P(\Delta W^*)$  is maximize when  $\Delta W_i^* = \kappa * \frac{\partial P}{\partial W_i}$ . Since  $P(\Delta W^*) = P(\Delta W_*)$  in  
345 the maximum, this proves the point.